

AI Forward Deployed Associate

Take-home case

Time: 3 to 4 hours of work, inside a 48-hour window. Use any AI tool you want, and tell us which ones you used and how. We are evaluating your judgment, not your unaided typing speed.

The dataset for this exercise, BrightChamps_FDA_Case_Dataset.csv, is in this folder. It holds 5,000 anonymised leads covering two months, with the following fields: lead id, lead source, geography, parent timezone, created timestamp, demo scheduled timestamp, rep assigned, rep shift, number of follow-up attempts, demo joined (Y/N), demo completed (Y/N), converted (Y/N).

Assume average revenue per converted customer of ₹60,000 and a blended marketing cost of ₹900 per lead.

DELIVER FOUR THINGS

- **1. The leak.** Where does this funnel lose the most money, and how much, in rupees per month? Show your working. We care more about how you sized it than about the exact figure.
- **2. The lever.** One intervention that could be live within two weeks with no engineering sprint. Tell us why you chose it over the alternatives you considered and rejected. The rejected options matter to us as much as the chosen one.
- **3. The build.** A working prototype of one component. An agent and its prompt, a script, a sheet with automation, a scoring model, whatever fits. It must actually run. Send it in a form we can execute or watch.
- **4. The memo.** One page maximum. The lever, the rupee number, the biggest adoption risk, how you would measure whether it worked, and the baseline you would register before starting.

RULES

Do not send a deck. A tight three-hour submission beats a sprawling ten-hour one.

If you make an assumption, write it down rather than asking us. The role works the same way.

Send everything as a single email with the memo as a PDF and the build as a file or a link we can open.

HOW TO SUBMIT

- Reply to the email that sent you this, as a single message, within 48 hours of receiving it.
- Attach the memo as a PDF. Send the build as a file, a repo link, or a live URL we can open without installing anything.
- Name the AI tools you used and what you used each one for. Using them is expected, not a mark against you.
- If anything is ambiguous, make a call, write down the assumption, and carry on. Do not wait on us to clarify.

How we will score it: problem framing and quantification (25%), lever selection and trade-off reasoning (25%), whether the build actually runs and could be owned by someone else (25%), adoption and measurement thinking (15%), written clarity (10%).